

**Name: Ritesh Chaudhary**

**Id: 2438464**

**Workshop week-7**

Ritesh\_Chaudhary\_2438464.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

RAM Disk Gemini

[1] import pandas as pd  
from matplotlib import pyplot as py  
from sklearn.model\_selection import train\_test\_split  
from sklearn.linear\_model import LogisticRegression  
from sklearn.metrics import accuracy\_score

[2] from google.colab import drive  
drive.mount('/content/drive')

Mounted at /content/drive

[3] df=pd.read\_csv("/content/drive/myDrive/Concept of AI Technology/creditcard.csv")  
df.head()

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153

5 rows x 31 columns

[4] df.shape

2s completed at 9:47 PM

```
Code | Text
✓ 0s [4] df.shape
(284807, 31)

✓ 0s df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Time    284807 non-null float64
1   V1      284807 non-null float64
2   V2      284807 non-null float64
3   V3      284807 non-null float64
4   V4      284807 non-null float64
5   V5      284807 non-null float64
6   V6      284807 non-null float64
7   V7      284807 non-null float64
8   V8      284807 non-null float64
9   V9      284807 non-null float64
10  V10     284807 non-null float64
11  V11     284807 non-null float64
12  V12     284807 non-null float64
13  V13     284807 non-null float64
14  V14     284807 non-null float64
15  V15     284807 non-null float64
16  V16     284807 non-null float64
17  V17     284807 non-null float64
18  V18     284807 non-null float64
19  V19     284807 non-null float64
20  V20     284807 non-null float64
21  V21     284807 non-null float64
22  V22     284807 non-null float64
23  V23     284807 non-null float64
24  V24     284807 non-null float64
25  V25     284807 non-null float64
26  V26     284807 non-null float64
27  V27     284807 non-null float64
28  V28     284807 non-null float64
29  V29     284807 non-null float64
30  V30     284807 non-null float64
31  V31     284807 non-null float64

✓ 2s completed
```

✓  
0s

[5]

30 Class 284807 non-null int64  
dtypes: float64(30), int64(1)  
memory usage: 67.4 MB

✓  
0s



df.isnull().sum()



time

0

V1 0

V2 0

V3 0

V4 0

V5 0

V6 0

V7 0

V8 0

V9 0

V10 0

V11 0

V12 0

V13 0

V14 0

V15 0

✓ 0s [7] df["Class"].value\_counts()

⇒

	count
Class	
0	284315
1	492

dtype: int64


✓ 0s [8] legit=df[df.Class==0]  
fraud=df[df.Class==1]

✓ 0s [9] legit.shape

⇒ (284315, 31)

✓ 0s [10] fraud.shape

⇒ (492, 31)

✓ 0s  legit.Amount.describe()

⇒

	Amount
count	284315.000000

```
[11] count    284315.000000
     mean      88.291022
     std       250.105092
     min         0.000000
     25%        5.650000
     50%       22.000000
     75%       77.050000
     max     25691.160000

dtype: float64
```

```
df.groupby('Class').mean()
```

```

      Time      V1      V2      V3      V4      V5      V6      V7      V8      V9      ...      V20      V21      V22      V23      V24      V25      V26
Class
0    94838.202258  0.008258 -0.006271  0.012171 -0.007860  0.005453  0.002419  0.009637 -0.000987  0.004467  ... -0.000844 -0.001235 -0.000024  0.000070  0.000182 -0.000072 -0.000081
1    80746.806911 -4.771948  3.623778 -7.033281  4.542029 -3.151225 -1.397737 -5.568731  0.570636 -2.581123  ...  0.372319  0.713588  0.014049 -0.040308 -0.105130  0.041449  0.051641
2 rows x 30 columns

```

```
[13] legit_sample= legit.sample(n=492) # Randomly selects 492 legitimate transactions (equal to the number of fraud cases).taking fraud ttransaction n=492 as it is fewer than legit so usir
```

```
[13] legit_sample= legit.sample(n=492) # Randomly selects 492 legitimate transactions (equal to the number of fraud cases).taking fraud ttransaction n=492 as it is fewer than legit so usir
```

```
[14] df_new= pd.concat([legit_sample, fraud], axis=0) #to create 50:50 ratio to balance the both legit and fraid dataset
```

```
[15] df_new.head()
```

```

      Time      V1      V2      V3      V4      V5      V6      V7      V8      V9      ...      V21      V22      V23      V24      V25      V26      V27
117742  74791.0 -0.057017  0.822460 -0.964387 -2.975748  1.373410 -1.818082  2.231105 -0.984326  0.441506  ...  0.293667  1.219066 -0.885084 -0.458677  1.385389 -0.254337 -0.272580
118627  75153.0  1.587825 -0.673364 -0.138171 -1.355835 -1.031468 -1.333353 -0.402736 -0.469623 -2.331611  ... -0.170811 -0.014042 -0.080359  0.414419  0.717290 -0.084751  0.005270
71509   54324.0 -2.726869 -2.949170  1.467949 -1.177913  1.158959 -0.115594 -0.395728  0.616507 -1.229687  ...  0.526731  0.212595  1.042204 -0.751253  0.216555 -0.318004 -0.107723
221716  142673.0 -0.868714  1.832264 -1.309321 -0.998193  1.262608  0.358677  0.072943 -2.356551  0.139450  ...  2.009316 -1.388955  0.441575 -0.274262 -0.814367  0.051324  0.420573
232379  147147.0  1.962747 -0.113934 -0.785193  0.637474 -0.432134 -0.682634 -0.528637 -0.036237  1.275986  ...  0.231150  0.885764  0.133406  1.077407 -0.057424 -0.197995  0.048019
5 rows x 31 columns

```

```
df_new['Class'].value_counts()
```

```

      count
Class
0         492
1         492

dtype: int64

```

```
[17] X= df_new.drop(columns='Class', axis=1) # will drop class column from dataset
     y= df_new['Class'] #since class is target variable it will contain both legit and fraud data corresponding to each row
```

```
[18] X_train, X_test, Y_train, Y_test=train_test_split(X,y,test_size=0.2,stratify=y, random_state=2) #stratify y to preserve the distribution of the target variable (y) across the training
```

```
[19] print(X_train.shape, Y_train.shape, X_test.shape, Y_test.shape)
```

```
(787, 30) (787, 1) (197, 30) (197, 1)
```

```
[20] lm = LogisticRegression()
```

```
✓ [21] lm.fit(X_train,Y_train)
0s
↳ /usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
n_iter_i = _check_optimize_result(
  ▾ LogisticRegression ⓘ ⓘ
  LogisticRegression()

✓ [22] X_train_predict=lm.predict(X_train)
0s
train_accuracy=accuracy_score(X_train_predict,Y_train)

✓ [23] print("accuracy:",train_accuracy)
0s
↳ accuracy: 0.9479034307496823
```

```
✓ [21] lm.fit(X_train,Y_train)
0s
↳ /usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
n_iter_i = _check_optimize_result(
  ▾ LogisticRegression ⓘ ⓘ
  LogisticRegression()

✓ [22] X_train_predict=lm.predict(X_train)
0s
train_accuracy=accuracy_score(X_train_predict,Y_train)

✓ [23] print("accuracy:",train_accuracy)
0s
↳ accuracy: 0.9479034307496823
```

✓ 0s [29] `model = LinearRegression()`  
`model.fit(X_train, y_train)`



LinearRegression ⓘ ?  
LinearRegression()

✓ 0s [30] `y_pred = model.predict(X_test)`

✓ 0s `mse = mean_squared_error(y_test, y_pred)`  
`print(f"Mean Squared Error: {mse:.2f}")`  
`print(f"Coefficient (slope): {model.coef_[0]:.2f}")`  
`print(f"Intercept: {model.intercept_:.2f}")`



Mean Squared Error: 3.68  
Coefficient (slope): -0.00  
Intercept: 0.16

✓ 2s `plt.scatter(X, y, color='blue', label="Original Data")`  
`plt.plot(X, model.predict(X), color='red', label="Regression Line")`  
`plt.xlabel("X (Feature)")`  
`plt.ylabel("y (Target)")`  
`plt.legend()`  
`plt.show()`



● Original Data



