

Name: Ritesh Chaudhary

Id: 2438464

## Workshop Week-2

Workshop2\_Ritesh\_Chaudhary\_2438464.ipynb ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

Name: Ritesh Chaudhary Id: 2438464

```
[13] #problem 1
#Dataset bank.csv
#1.
import pandas as pd
df=pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/bank .csv')
df.head()
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261	1	-1	0	unknown	no
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151	1	-1	0	unknown	no
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	76	1	-1	0	unknown	no
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	92	1	-1	0	unknown	no
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may	198	1	-1	0	unknown	no

```
#2.
# Check the info of the DataFrame
df_info = df.info()
print(df_info)

object_columns = df.select_dtypes(include=['object']).columns
print("\nColumns with dtype=object:")
print(object_columns)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   x        100 non-null    float64
1   y        100 non-null    float64
dtypes: float64(2)
memory usage: 1.7 KB
None

Columns with dtype=object:
Index([], dtype='object')
```

✓  
0s



#2(b)

```
import pandas as pd

# Load the dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/bank .csv')

# Get info of the DataFrame
df.info()

# Identify columns with dtype 'object'
object_columns = df.select_dtypes(include=['object']).columns

for col in object_columns:
    print(f"Unique values in '{col}':")
    print(df[col].unique())
    print("\n")
```



3	education	45211	non-null	object
4	default	45211	non-null	object
5	balance	45211	non-null	int64
6	housing	45211	non-null	object
7	loan	45211	non-null	object
8	contact	45211	non-null	object
9	day	45211	non-null	int64
10	month	45211	non-null	object
11	duration	45211	non-null	int64
12	campaign	45211	non-null	int64
13	pdays	45211	non-null	int64
14	previous	45211	non-null	int64
15	poutcome	45211	non-null	object
16	y	45211	non-null	object

```
✓ 0s [▶] 15 poutcome 45211 non-null object
16 y 45211 non-null object
[↕] dtypes: int64(7), object(10)
memory usage: 5.9+ MB
Unique values in 'job':
['management' 'technician' 'entrepreneur' 'blue-collar' 'unknown'
'retired' 'admin.' 'services' 'self-employed' 'unemployed' 'housemaid'
'student']

Unique values in 'marital':
['married' 'single' 'divorced']

Unique values in 'education':
['tertiary' 'secondary' 'unknown' 'primary']

Unique values in 'default':
['no' 'yes']

Unique values in 'housing':
['yes' 'no']

Unique values in 'loan':
['no' 'yes']

Unique values in 'contact':
['unknown' 'cellular' 'telephone']
```

```
✓ 0s [▶] [ 'unknown' 'cellular' 'telephone' ]
[↕] Unique values in 'month':
['may' 'jun' 'jul' 'aug' 'oct' 'nov' 'dec' 'jan' 'feb' 'mar' 'apr' 'sep']

Unique values in 'poutcome':
['unknown' 'failure' 'other' 'success']

Unique values in 'y':
['no' 'yes']
```

+ Code + Text

✓  
0s

```
#2(c)
import pandas as pd

# Load the dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/bank .csv')

# Check for the total number of null values in each column
null_values = df.isnull().sum()

# Print the result
print(null_values)
```

```
↕ age      0
   job      0
   marital  0
   education  0
   default  0
   balance  0
   housing  0
   loan     0
   contact  0
   day      0
   month    0
   duration  0
   campaign  0
   pdays    0
   previous  0
   poutcome  0
   y        0
   dtype: int64
```

+ Code + Text

```
[16] #3
import pandas as pd

# Load the dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/bank .csv')

# Drop columns with dtype 'object' and store in a new DataFrame
df_numeric = df.select_dtypes(exclude=['object'])

# Save the new DataFrame to a CSV file
df_numeric.to_csv('banknumericdata.csv', index=False)

# Optional: print the first few rows of the new DataFrame to verify
print(df_numeric.head())
```

	age	balance	day	duration	campaign	pdays	previous
0	58	2143	5	261	1	-1	0
1	44	29	5	151	1	-1	0
2	33	2	5	76	1	-1	0
3	47	1506	5	92	1	-1	0
4	33	1	5	198	1	-1	0

+ Code + Text

```
[17] #4
import pandas as pd

# Read the 'banknumericdata.csv' file
df_numeric = pd.read_csv('banknumericdata.csv')

# Get the summary statistics for the numeric columns
summary_statistics = df_numeric.describe()

# Print the summary statistics
print(summary_statistics)
```

	age	balance	day	duration	campaign	\
count	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000	
mean	40.936210	1362.272058	15.806419	258.163080	2.763841	
std	10.618762	3044.765829	8.322476	257.527812	3.098021	
min	18.000000	-8019.000000	1.000000	0.000000	1.000000	
25%	33.000000	72.000000	8.000000	103.000000	1.000000	
50%	39.000000	448.000000	16.000000	180.000000	2.000000	
75%	48.000000	1428.000000	21.000000	319.000000	3.000000	
max	95.000000	102127.000000	31.000000	4918.000000	63.000000	

  

	pdays	previous
count	45211.000000	45211.000000
mean	40.197828	0.580323
std	100.128746	2.303441
min	-1.000000	0.000000
25%	-1.000000	0.000000
50%	-1.000000	0.000000
75%	-1.000000	0.000000
max	871.000000	275.000000

+ Code + Text

```
[18] #Problem 2 Data imputations
#1.
import pandas as pd

# Load the 'medical student' dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/medical_students_dataset.csv')

# Print the first few rows to verify
print(df.head())
```

	Student ID	Age	Gender	Height	Weight	Blood Type	BMI	\
0	1.0	18.0	Female	161.777924	72.354947	O	27.645835	
1	2.0	NaN	Male	152.069157	47.630941	B	NaN	
2	3.0	32.0	Female	182.537664	55.741083	A	16.729017	
3	NaN	30.0	Male	182.112867	63.332207	B	19.096042	
4	5.0	23.0	Female	NaN	46.234173	O	NaN	

  

	Temperature	Heart Rate	Blood Pressure	Cholesterol	Diabetes	Smoking
0	NaN	95.0	109.0	203.0	No	NaN
1	98.714977	93.0	104.0	163.0	No	No
2	98.260293	76.0	130.0	216.0	Yes	No
3	98.839605	99.0	112.0	141.0	No	Yes
4	98.480008	95.0	NaN	231.0	No	No

+ Code + Text

```
#2
import pandas as pd

# Load the 'medical student' dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/medical_students_dataset.csv')

# Check the info of the DataFrame
df.info()

# Identify columns with missing (null) values
missing_values = df.isnull().sum()

# Print columns with missing values
print("Columns with missing values:")
print(missing_values[missing_values > 0])
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 13 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   Student ID      180000 non-null float64
 1   Age             180000 non-null float64
 2   Gender          180000 non-null object
 3   Height          180000 non-null float64
 4   Weight          180000 non-null float64
 5   Blood Type     180000 non-null object
 6   BMI             180000 non-null float64
 7   Temperature    180000 non-null float64
 8   Heart Rate     180000 non-null float64
 9   Blood Pressure 180000 non-null float64
10   Cholesterol     180000 non-null float64
```

```
[20] 10 Cholesterol      180000 non-null float64
      11 Diabetes      180000 non-null object
      12 Smoking       180000 non-null object
dtypes: float64(9), object(4)
memory usage: 19.8+ MB
Columns with missing values:
Student ID      20000
Age             20000
Gender          20000
Height         20000
Weight         20000
Blood Type     20000
BMI            20000
Temperature    20000
Heart Rate     20000
Blood Pressure 20000
Cholesterol    20000
Diabetes       20000
Smoking        20000
dtype: int64
```

```

[21] #3
import pandas as pd

# Load the 'medical student' dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/medical_students_dataset.csv')

# Check for missing values
missing_values = df.isnull().sum()
print("Columns with missing values:\n", missing_values[missing_values > 0])

if 'age' in df.columns:
    # If 'age' is skewed, use median to fill missing values
    df['age'] = df['age'].fillna(df['age'].median())
    print("Missing values in 'age' filled using median.")

if 'gender' in df.columns:
    # Fill missing values in 'gender' with the most frequent value (mode)
    df['gender'] = df['gender'].fillna(df['gender'].mode()[0])
    print("Missing values in 'gender' filled using mode.")

if 'timestamp' in df.columns:
    df['timestamp'] = df['timestamp'].fillna(method='ffill')
    print("Missing values in 'timestamp' filled using forward fill.")

# Check if all missing values have been filled
print("\nMissing values after imputation:\n", df.isnull().sum())

```

Columns with missing values:

Student ID	20000
------------	-------

✓ 0s completed at 8:06 PM



0s



Columns with missing values:



Student ID	20000
Age	20000
Gender	20000
Height	20000
Weight	20000
Blood Type	20000
BMI	20000
Temperature	20000
Heart Rate	20000
Blood Pressure	20000
Cholesterol	20000
Diabetes	20000
Smoking	20000
dtype: int64	

Missing values after imputation:

Student ID	20000
Age	20000
Gender	20000
Height	20000
Weight	20000
Blood Type	20000
BMI	20000
Temperature	20000
Heart Rate	20000
Blood Pressure	20000
Cholesterol	20000
Diabetes	20000
Smoking	20000
dtype: int64	

```

#4
import pandas as pd

# Load the 'medical student' dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/medical_students_dataset.csv')

# Check for duplicate rows
duplicate_rows = df.duplicated().sum()

# Print the number of duplicate rows
print(f"Number of duplicate rows: {duplicate_rows}")

# If there are duplicates, remove them
if duplicate_rows > 0:
    # Remove duplicate rows
    df_cleaned = df.drop_duplicates()

    # Print the number of rows after removing duplicates
    print(f"Number of rows after removing duplicates: {df_cleaned.shape[0]}")
else:
    print("No duplicates found.")

# Optionally: Save the cleaned DataFrame to a new CSV file
df_cleaned.to_csv('medical_student_cleaned.csv', index=False)

# Check the first few rows of the cleaned DataFrame
print("\nFirst few rows after cleaning:")
print(df_cleaned.head())

```

```

Number of duplicate rows: 7644
Number of rows after removing duplicates: 192356

```

```

Number of duplicate rows: 7644
Number of rows after removing duplicates: 192356

```

First few rows after cleaning:

	Student ID	Age	Gender	Height	Weight	Blood Type	BMI	\
0	1.0	18.0	Female	161.777924	72.354947	O	27.645835	
1	2.0	NaN	Male	152.069157	47.630941	B	NaN	
2	3.0	32.0	Female	182.537664	55.741083	A	16.729017	
3	NaN	30.0	Male	182.112867	63.332207	B	19.096042	
4	5.0	23.0	Female	NaN	46.234173	O	NaN	

	Temperature	Heart Rate	Blood Pressure	Cholesterol	Diabetes	Smoking
0	NaN	95.0	109.0	203.0	No	NaN
1	98.714977	93.0	104.0	163.0	No	No
2	98.260293	76.0	130.0	216.0	Yes	No
3	98.839605	99.0	112.0	141.0	No	Yes
4	98.480008	95.0	NaN	231.0	No	No



```
#Exercise 3.2
#Dataset name 'titanic.csv'
#Problem 1
import pandas as pd

# Load the Titanic dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/Titanic-Dataset.csv')

# Subset the DataFrame to include only relevant columns
df_subset = df[['Name', 'Pclass', 'Sex', 'Age', 'Fare', 'Survived']]

# Retain only rows where Pclass is equal to 1 (first-class passengers)
df_first_class = df_subset[df_subset['Pclass'] == 1]

# Calculate the mean, median, maximum, and minimum of the 'Fare' column
mean_fare = df_first_class['Fare'].mean()
median_fare = df_first_class['Fare'].median()
max_fare = df_first_class['Fare'].max()
min_fare = df_first_class['Fare'].min()

# Print the results
print(f"Mean Fare: {mean_fare}")
print(f"Median Fare: {median_fare}")
print(f"Maximum Fare: {max_fare}")
print(f"Minimum Fare: {min_fare}")
```



```
Mean Fare: 84.1546875
Median Fare: 60.287499999999994
Maximum Fare: 512.3292
Minimum Fare: 0.0
```

0s



### #Problem 2

```
import pandas as pd

# Load the Titanic dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/Titanic-Dataset.csv')

# Subset the DataFrame to include only relevant columns
df_subset = df[['Name', 'Pclass', 'Sex', 'Age', 'Fare', 'Survived']]

df_first_class = df_subset[df_subset['Pclass'] == 1]

# Check the number of null values in the 'Age' column
null_age_count = df_first_class['Age'].isnull().sum()

# Print the number of null values
print(f"Number of null values in 'Age' column: {null_age_count}")

# Drop rows with missing values in the 'Age' column
df_first_class_cleaned = df_first_class.dropna(subset=['Age'])

# Verify the changes
print(f"Number of rows after dropping null values: {df_first_class_cleaned.shape[0]}")
```



Number of null values in 'Age' column: 30  
Number of rows after dropping null values: 186

0s



### #Problem 3

```
import pandas as pd

# Load the Titanic dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/Titanic-Dataset.csv')

# Check the unique values in the 'Embarked' column
print("Unique values in 'Embarked' column:")
print(df['Embarked'].unique())

# Check for missing values in 'Embarked'
missing_embarked = df['Embarked'].isnull().sum()
print(f"\nNumber of missing values in 'Embarked' column: {missing_embarked}")

# Handle missing values - We can either drop or fill with the mode (most frequent value)
# Let's fill missing 'Embarked' values with the mode (most frequent port)
df['Embarked'] = df['Embarked'].fillna(df['Embarked'].mode()[0])

# Check the frequency of each embarkation port after filling missing values
embarked_freq = df['Embarked'].value_counts()
print("\nFrequency of each embarkation port:")
print(embarked_freq)

# Display the first few rows of the DataFrame after handling the 'Embarked' column
print("\nFirst few rows after handling 'Embarked' column:")
print(df[['Name', 'Pclass', 'Embarked']].head())
```



Unique values in 'Embarked' column:  
['S' 'C' 'Q' nan]

Unique values in 'Embarked' column:  
['S' 'C' 'Q' nan]

Number of missing values in 'Embarked' column: 2

Frequency of each embarkation port:

Embarked

S 646

C 168

Q 77

Name: count, dtype: int64

First few rows after handling 'Embarked' column:

	Name	Pclass	Embarked
0	Braund, Mr. Owen Harris	3	S
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	1	C
2	Heikkinen, Miss. Laina	3	S
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	1	S
4	Allen, Mr. William Henry	3	S

[27] #1.

```
import pandas as pd
```

```
# Load the Titanic dataset
```

```
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/Titanic-Dataset.csv')
```

```
# Use one-hot encoding to convert 'Embarked' into separate binary columns
```

```
df_encoded = pd.get_dummies(df['Embarked'], prefix='Embarked')
```

```
# Print the first few rows of the new one-hot encoded columns
```

```
print(df_encoded.head())
```

	Embarked_C	Embarked_Q	Embarked_S
0	False	False	True
1	True	False	False
2	False	False	True
3	False	False	True
4	False	False	True

```
#2.

import pandas as pd

# Load the Titanic dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/Titanic-Dataset.csv')

# Check the unique values in the 'Embarked' column to see if it matches 'C', 'Q', 'S'
print(df['Embarked'].unique())

# Use one-hot encoding to convert 'Embarked' into separate binary columns
df_encoded = pd.get_dummies(df['Embarked'], prefix='Embarked')

# Print the column names of the encoded DataFrame to ensure the columns were created
print(df_encoded.columns)

# Add the new one-hot encoded columns to the original DataFrame
df = pd.concat([df, df_encoded], axis=1)

# Print the first few rows of the modified DataFrame to verify the changes
print(df[['Name', 'Pclass'] + list(df_encoded.columns)].head())
```

```
['S' 'C' 'Q' nan]
Index(['Embarked_C', 'Embarked_Q', 'Embarked_S'], dtype='object')
   Name  Pclass  Embarked_C \
0  Braund, Mr. Owen Harris      3      False
1  Cumings, Mrs. John Bradley (Florence Briggs Th...      1       True
2  Heikkinen, Miss. Laina      3      False
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)      1      False
4  Allen, Mr. William Henry      3      False

   Embarked_Q  Embarked_S
```

```
['S' 'C' 'Q' nan]
Index(['Embarked_C', 'Embarked_Q', 'Embarked_S'], dtype='object')
   Name  Pclass  Embarked_C \
0  Braund, Mr. Owen Harris      3      False
1  Cumings, Mrs. John Bradley (Florence Briggs Th...      1       True
2  Heikkinen, Miss. Laina      3      False
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)      1      False
4  Allen, Mr. William Henry      3      False

   Embarked_Q  Embarked_S
0      False       True
1      False      False
2      False       True
3      False       True
4      False       True
```



```
#3.
import pandas as pd

# Load the Titanic dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/Titanic-Dataset.csv')

# Check the unique values in the 'Embarked' column to see if it matches 'C', 'Q', 'S'
print(df['Embarked'].unique())

# Use one-hot encoding to convert 'Embarked' into separate binary columns
df_encoded = pd.get_dummies(df['Embarked'], prefix='Embarked')

# Add the new one-hot encoded columns to the original DataFrame
df = pd.concat([df, df_encoded], axis=1)

# Drop the original 'Embarked' column
df = df.drop(columns=['Embarked'])

# Print the first few rows of the modified DataFrame to verify the changes
print(df[['Name', 'Pclass'] + list(df_encoded.columns)].head())
```



['S' 'C' 'Q' nan]

	Name	Pclass	Embarked_C	\
0	Braund, Mr. Owen Harris	3	False	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	1	True	
2	Heikkinen, Miss. Laina	3	False	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	1	False	
4	Allen, Mr. William Henry	3	False	

  

	Embarked_Q	Embarked_S
0	False	True
1	False	False



['S' 'C' 'Q' nan]

	Name	Pclass	Embarked_C	\
0	Braund, Mr. Owen Harris	3	False	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	1	True	
2	Heikkinen, Miss. Laina	3	False	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	1	False	
4	Allen, Mr. William Henry	3	False	

  

	Embarked_Q	Embarked_S
0	False	True
1	False	False
2	False	True
3	False	True
4	False	True

```

import pandas as pd

# Load the Titanic dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/Titanic-Dataset.csv')

# Check the unique values in the 'Embarked' column to see if it matches 'C', 'Q', 'S'
print("Unique values in 'Embarked' column:", df['Embarked'].unique())

# Use one-hot encoding to convert 'Embarked' into separate binary columns
df_encoded = pd.get_dummies(df['Embarked'], prefix='Embarked')

# Print the columns of the one-hot encoded DataFrame to check what was created
print("Columns created by one-hot encoding:", df_encoded.columns)

# Add the new one-hot encoded columns to the original DataFrame
df = pd.concat([df, df_encoded], axis=1)

# Drop the original 'Embarked' column
df = df.drop(columns=['Embarked'])

# Print the first few rows of the modified DataFrame to verify the changes
print(df[['Name', 'Pclass'] + list(df_encoded.columns)].head())

```

```

Unique values in 'Embarked' column: ['S' 'C' 'Q' nan]
Columns created by one-hot encoding: Index(['Embarked_C', 'Embarked_Q', 'Embarked_S'], dtype='object')

```

	Name	Pclass	Embarked_C	\
0	Braund, Mr. Owen Harris	3	False	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	1	True	
2	Heikkinen, Miss. Laina	3	False	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	1	False	
4	Allen, Mr. William Henry	3	False	

```

#Problem 4
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load the Titanic dataset
df = pd.read_csv('/content/drive/MyDrive/Concept of Ai Technology/Titanic-Dataset.csv')

# Group by 'Sex' and calculate the mean of 'Survived' for each group
mean_survival_by_gender = df.groupby('Sex')['Survived'].mean()

# Print the mean survival rates by gender
print(mean_survival_by_gender)

# Create a bar plot to visualize the survival rates by gender
plt.figure(figsize=(8, 6))
sns.barplot(x=mean_survival_by_gender.index, y=mean_survival_by_gender.values, palette='viridis')

# Set the title and labels
plt.title('Mean Survival Rates by Gender')
plt.xlabel('Gender')
plt.ylabel('Mean Survival Rate')
plt.show()

```

```

Sex
female    0.742038
male      0.188908
Name: Survived, dtype: float64
<ipython-input-31-6cdc823fd9c9>:17: FutureWarning:

```

Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'x' variable to 'hue' and set 'legend=False' for the same effect.





